


Roleplaying probes across model depth

Same 371 paired scenarios | Test: 56 held-out scenarios, 112 completions



Stars: validation-selected blocks (Qwen 14, Llama 18; zero-based). C selected separately per layer on validation.

Frozen replay of identical completion texts; completion-token mean pooling. Depth normalization does not equate blocks.