


Text and length nearly separate report labels

Initial Llama insider-trading evaluation | 200 held-out trajectories from the same setting



Report-token mean activation (selected block 0), report-text TF-IDF, and report-token length; all use L2 logistic regression.

Dotted lines: constant-classifier baselines. Accuracy uses a 0.5 threshold; TF-IDF ranks well despite low thresholded accuracy.

Near-perfect ranking from text and length means report-probe performance alone does not isolate a deception mechanism.